

Networking II

Joe Rogers
Android 310

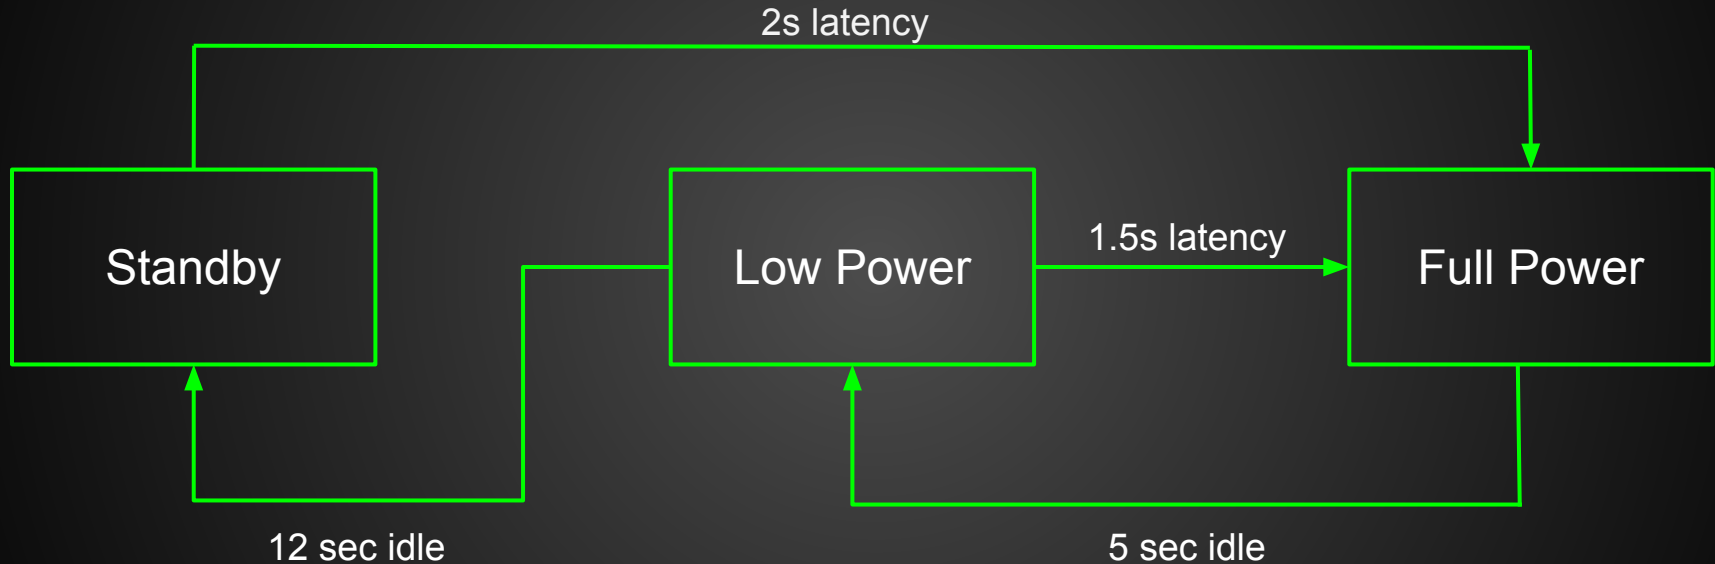
Why worry about networking

- Developers do not always worry about network impacts
 - Tend to develop on wifi and stable environment
 - Thinking of features and not the impact on data usage and battery life.
- However, as Android has progressed, the OS is fingering apps that consume large amounts of data and battery.

First instincts are not always right

- When thinking of network usage, most developer's first instincts are:
 - Fetch data when user opens the app
 - Fetch only what I need for the screen I'm showing.
- Seems reasonable. Developer limiting data use by only fetching data on an as needed basis.

Mobile Radio State Machine



AT&T 3G Network

Ad-hoc fetching

- Lets say app fetches data every 20 second for 1 sec based on user action 3 times. Ie nav to new screen, etc.
- Spends 18 sec at full power. (3 sec of transfer time + 15 sec idle).
- Spends 42 seconds at low power. (36 sec at low power idle + ~5 sec in transition to high power)
- Therefore over the course of a minute the radio was never idle or was there for less that 1 sec, depending on transition timing.

Bundled fetch

- Instead of 3 separate fetches, app fetches all data at once.
- Spends 8 sec at full power. (3 sec + 5 sec full power idle).
- Spends ~14 sec at low power (12 sec idle, plus 2 sec transition to full power)
- Spends ~38 sec idle for the minute.
- This is a notable improvement in battery life for the same amount of data.

Ad-hoc fetching (User perspective)

- Every time user transitions, data is being “fetched” possibly showing a loading indicator.
 - Great on wifi, or high speed lte. Likely very short wait times.
 - Not so good when you get to 3G/2G or a poorly performing network. Those waiting indicators start to go up.

Bundled fetch (User perspective)

- First loads app, may see a loading indicator especially on slow networks.
- Generally for the 2nd/3rd fetches in a bundled fetch, the user should not see a loading indicator as the data is already downloaded. This also improves the user's experience.

Prefetching data

- Essentially front loading data downloads when user starts the app.
- Knowing that a single fetch is going to power the radio for at least 20 seconds, fetching additional information to conserve battery.
- Another advantage is the user not waiting for certain data to download.

Prefetching data cont.

- Based on the radio state machine. Ideally you fetch data that has a 50% chance of being used during the user session.
- Put another way, using the state machine as a guide, in the current user session you could download data for ~6 seconds which would yield up to ~1 to 2MB of data.

Prefetching data cont.

- However, don't spend 6 seconds and 1-2 MB of downloads to just to fetch data if it is likely to be used less than 50% of time. It is still better to stop earlier.

Real World Example: News app

- Bad App
 - Some news apps fetch only the headlines for the category being shown.
 - Fetch the stories as read and thumbnails the first time they become visible.
 - Results in a lot of random network traffic. Especially if the user is slowly moving up and down the list maybe after reading one story.

Real World Example: News App

- Good app
 - Fetches headlines/thumbs for first category. Then continues to fetch full stories for first category, and headlines/thumbs for other categories as time permits.
 - Even better is app tracks categories user prefers to open and “prefetches” those categories.
 - Alternatively the server tracks stories users tend to open and move those to front of list.

Real World Example: Music app

- While album playing one instinct is to prefetch all of the tracks. However, if user stops after current song, a lot went wasted.
- Better solution is to fetch next track in playlist as user has a 50% chance to listen to it.

The Weather App

- The intent service technique of fetching the current conditions and forecast data in one shot is an example of prefetching data the user may want to see.
 - If the user opens the app likely want to see current conditions 100% of the time.
 - May click to a forecast item 50% of the time especially if its especially good or bad weather.

Batching data

- Another thing apps can do is defer transfers until the next “must” transfer.
 - For example analytic data.
 - Instead of sending data to server as collected, the app should persist the data and then send it the next time it “fetches” data for the user.
 - This may require a little hand-holding with management that wants to see the analytic data ASAP but is much better for battery life.

Background fetching

- Depending on the nature of the app, it may be worth prefetching data in the background in anticipation of a user session.
- However, the risk of the user not seeing the data is much higher so the threshold for performing this should also be higher and less frequent.

Background fetching cont.

- Some ways to increase threshold are:
 - Only fetching if on wifi
 - Only fetching if charging
 - Fetching infrequently based on knowledge of the data or user patterns.

Background fetching example

- Sports app
 - Could fetch upcoming schedule once a week for favorite teams and leagues.
 - If a favorite team is “playing” fetch latest score (or scores for league) every 5-10 min in anticipation of the user checking the score of the league.
 - Or fetch when “notification” of game event arrives anticipating the user opening the app to check progress.

Weather app example

- Fetch current conditions every few hours
 - The “free” API states it may “fail” and to check back in 10 min is not a good experience for the user.
Having data 2-4 hours old is better than days old...
- Fetch forecasts once a day.
 - One doesn’t change much and if “fail” to get data likely have a few days in reserve.
- Ideally user chooses to sync/not sync.

Solving the problem

Reduce Data transfer size

- **Problem:** Uncompressed data is significantly more expensive to fetch with respect to radio/latency.
- **Solution:** Ensure server is gzip enabled.
URLConnection requests gzip data by default.
 - Compressed Json/XML is roughly 10x smaller than uncompressed. 50kb vs 500kb.
 - Actually proved this to get server team to enable compression on a project.
 - Will use less battery decompressing vs downloading

Reducing data fetches

- Bundle your fetches
 - As shown, fetching data that will be used 50% of time (up to 6 sec/1-2mb) will be better than than fetching data piecemeal.
- Use caching
 - Use response cache, and/or local cache (database) to reduce data fetches.
 - If data “could” be cached, work with server team to ensure headers are configured to avoid network.

Reducing data fetches cont.

- Know your data
 - If data changes no faster than an hour, ensure you don't fetch every 5 minutes. If you use a database, query for last fetch time.
 - If background fetches are important, ensure you don't fetch too frequently or the cost/benefit will be reduced.

Know your network

- Adjust download behavior depending on network type.
 - Download more on wifi
 - Possibly avoid certain data on slower networks (2G).
For example on news app, avoid thumbnail images on alternate news categories etc.
 - Also if on slow network, avoid any background operations.

Monitor connectivity

- In last weeks sample if no network, the app did a “backoff/retry” solution.
- A better solution is to use broadcast receiver to monitor the `CONNECTIVITY_ACTION` defined by the connectivity manager. When network is restored, the BR would start the intent service if app still active.

Background Techniques

From best to worst

Use Google Cloud Messaging

- Server sends a GCM message when the data changes. Client not “polling”
 - Requires server and client setup.
 - When client receives message, it schedules a job to download updated data.
 - Example: new email arrives. Gmail server send GCM message to client. Client knows to sync to fetch new messages and notify the user.

Use Job Scheduler

- Job scheduler lets client be woken when certain conditions met such as time, network enabled, charging, etc.
 - Client not woken if condition not met.
 - Client notified if conditions change to stop work
 - OS will “bundle” jobs together. If 3 jobs need network, then OS schedules to use radio together.
 - New in Android 5.0 (but no compatibility)
 - Will persist across boots if desired.

Alarm Manager

- Used as job scheduler prior to 5.0
- Can be used to schedule repeating, exact and inexact events.
- Inexact alarms will not go off precisely every 15 min to allow different alarms to be fire at once.
- No idea if app only one requiring network and network may not be available, etc.
- Must manually reschedule alarm after boot.

Resources

Resources

- [Transferring data minimizing battery impact](#)
- [Understanding the cell radio](#) (video)
- [Job Scheduler](#)
- [Connectivity Manager](#)
- [Alarm Manager](#)